

## 0.1 Diagnostic Plots for Linear Models with R

Plot Diagnostics for an `lm` Object

Six plots (selectable by `which`) are currently available:

1. a plot of residuals against fitted values,
2. a Scale-Location plot of  $\sqrt{|\text{residuals}|}$  against fitted values,
3. a Normal Q-Q plot,
4. a plot of Cook's distances versus row labels,
5. a plot of residuals against leverages,
6. a plot of Cook's distances against leverage/(1-leverage).

By default, the first three and 5 are provided.

### Residuals plots

`lme` allows to plot the residuals in the following ways:

```
res_lme=residuals(model_lme)
plot(res_lme)
qqnorm(res_lme)
qqline(res_lme)
plot(model_lme)
```

When the `plot` function calls the model object, the residual plot is produced.

```
plot(JS.roy1, which=c(1) )
```

LME models assume that the residuals of the model are normally distributed. A Normal probability plot can be constructed to check this assumption. Commonly used R commands can be used to construct the plot.

```
qqnorm(resid(JS.roy1),pch="*",col="red")  
qqline(resid(JS.roy1),col="blue")
```

```
table(dat$method[1:255])  
##  
##    J    S  
## 255    0  
table(dat$method[256:510])  
##  
##    J    S  
##    0 255
```

```
library(predictMeans)  
CookD(model, group=method, plot=TRUE, idn=5, newwd=FALSE)
```

```
> shapiro.test(resid(JS.roy1)[256:510])
```

Shapiro-Wilk normality test

data: resid(JS.roy1)[256:510]

W = 0.9395, p-value = 9.503e-09

```
plot(roy.NLME, resid(., type = "p") ~ fitted(.) | method,  
abline = 0, id=.05)
```

```
library(predictMeans)  
CookD(model, group=method, plot=TRUE, idn=5, newwd=FALSE)
```

```
blood.red <- blood[!(blood$subject %in% c(68,78,80)),]
```

```
dim(blood.red)
```

```
# 27 observations should be removed.
```

```
blood.NLME.red <- lme(BP ~ method-1 , random=~1|subject,data = blood.red)
```

```
plot(blood.NLME.red, resid(., type = "p") ~ fitted(.) | method, abline = 0, id=.05)
```

```
> shapiro.test(resid(JS.roy1)[1:255])
```

Shapiro-Wilk normality test

data: resid(JS.roy1)[1:255]

W = 0.9931, p-value = 0.2852

```
> shapiro.test(resid(JS.roy1)[256:510])
```

Shapiro-Wilk normality test

data: resid(JS.roy1)[256:510]

W = 0.9395, p-value = 9.503e-09

```
data.frame( response = resid(JS.ARoy20091, type = "response"),  
  pearson = resid(JS.ARoy20091, type = "pearson"),  
  normalized = resid(JS.ARoy20091, type = "normalized") )
```

	response	pearson	normalized
1	-4.65805902	-0.761587227	-0.7615872269
2	-0.88701342	-0.145025661	0.0776238081
3	-5.16580898	-0.844603753	-0.8446037530

```

4      2.29041830  0.374480726  0.6450898404
5      7.87508366  1.287567009  1.2875670086
6     -6.57048659 -1.074266908 -1.5090772378
.....

```

For the  $J$  observations, the variance is 6.116252 whereas for the  $S$  observations, the denominator is 9.118144. (with the expected ratio of 1.490806)

```

> pearson %>%
+   as.numeric %>%
+   matrix(nrow=85) %>%
+   round(4)
[,1]    [,2]    [,3]    [,4]    [,5]    [,6]
[1,] -0.7616  0.2194  0.3829 -0.2983  0.3597 -0.0790
[2,] -0.1450  0.1820 -0.1450 -0.5014  0.1567  0.2663
[3,] -0.8446  0.4634  0.1364 -0.1630 -0.2727  0.1660
[4,]  0.3745 -0.2795 -0.2795 -0.2658 -0.2658  0.6115
[5,]  1.2876 -0.6744 -0.6744  0.8935 -0.0935 -0.8612
[6,] -1.0743  1.8687 -0.7473 -0.0383  0.2908 -0.3673
.....

```

We can plot the residuals against the fitted values, to assess the assumption of constant variance.

```
# standardized residuals versus fitted values  
plot(JS.ARoy20091, resid(., type = "pearson") ~ fitted(.) ,  
      abline = 0, id = 0.05)
```

```
par(mfrow=c(1,2))  
qqnorm((resid(JS.ARoy20091)[1:255]),  
        pch="*",col="red",  
        ylim=c(-40,40),  
        main="Method J")  
qqline(resid(JS.ARoy20091)[1:255],col="blue")  
qqnorm((resid(JS.ARoy20091)[256:510]),  
        pch="*",col="red",  
        ylim=c(-40,40),  
        main="Method S")  
qqline(resid(JS.ARoy20091)[256:510],col="blue")  
par(mfrow=c(1,1))
```

## Residuals plots

When the `plot` function calls the model object, the residual plot is produced.

```
plot(JS.roy1, which=c(1) )
```

LME models assume that the residuals of the model are normally distributed. A Normal probability plot can be constructed to check this assumption. Commonly used R commands can be used to construct the plot.

```
qqnorm(resid(JS.roy1),pch="*",col="red")
qqline(resid(JS.roy1),col="blue")
```

```
table(dat$method[1:255])
##
##    J    S
## 255    0
table(dat$method[256:510])
##
##    J    S
##    0 255
```

```
plot(roy.NLME, resid(., type = "p") ~ fitted(.) | method,
     abline = 0, id=.05)
```

```
data.frame( response = resid(JS.ARoy20091, type = "response"),
  pearson   = resid(JS.ARoy20091, type = "pearson"),
  normalized = resid(JS.ARoy20091, type = "normalized") )
```

	response	pearson	normalized
1	-4.65805902	-0.761587227	-0.7615872269
2	-0.88701342	-0.145025661	0.0776238081
3	-5.16580898	-0.844603753	-0.8446037530
4	2.29041830	0.374480726	0.6450898404
5	7.87508366	1.287567009	1.2875670086
6	-6.57048659	-1.074266908	-1.5090772378
.....			

For the  $J$  observations, the variance is 6.116252 whereas for the  $S$  observations, the denominator is 9.118144. (with the expected ratio of 1.490806)

```
> pearson %>%
+   as.numeric %>%
+   matrix(nrow=85) %>%
+   round(4)

[,1]    [,2]    [,3]    [,4]    [,5]    [,6]
[1,] -0.7616  0.2194  0.3829 -0.2983  0.3597 -0.0790
[2,] -0.1450  0.1820 -0.1450 -0.5014  0.1567  0.2663
[3,] -0.8446  0.4634  0.1364 -0.1630 -0.2727  0.1660
[4,]  0.3745 -0.2795 -0.2795 -0.2658 -0.2658  0.6115
[5,]  1.2876 -0.6744 -0.6744  0.8935 -0.0935 -0.8612
[6,] -1.0743  1.8687 -0.7473 -0.0383  0.2908 -0.3673
.....
```



We can plot the residuals against the fitted values, to assess the assumption of constant variance.

```
# standardized residuals versus fitted values
plot(JS.ARoy20091, resid(., type = "pearson") ~ fitted(.) ,
     abline = 0, id = 0.05)
```

```
par(mfrow=c(1,2))
qqnorm((resid(JS.ARoy20091)[1:255]),
      pch="*", col="red",
      ylim=c(-40,40),
      main="Method J")
qqline(resid(JS.ARoy20091)[1:255], col="blue")
qqnorm((resid(JS.ARoy20091)[256:510]),
      pch="*", col="red",
      ylim=c(-40,40),
      main="Method S")
qqline(resid(JS.ARoy20091)[256:510], col="blue")
par(mfrow=c(1,1))
```

This code will allow you to make QQ plots for each level of the random effects. LME models assume that not only the within-cluster residuals are normally distributed, but that each level of the random effects are as well. Depending on the model, you can vary the level from 0, 1, 2 and so on

```
qqnorm(JS.ARoy20091, ~ranef(.))
```

```
# qqnorm(JS.ARoy20091, ~ranef(.,levels=1))
```




Figure 1:

```
data.frame( response = resid(JS.roy1, type = "response"),
  pearson = resid(JS.roy1, type = "pearson"),
  normalized = resid(JS.roy1, type = "normalized") )
```

```
response    pearson    normalized
1    -4.65805902 -0.761587227 -0.7615872269
```

```

2    -0.88701342 -0.145025661  0.0776238081
3    -5.16580898 -0.844603753 -0.8446037530
4     2.29041830  0.374480726  0.6450898404
5     7.87508366  1.287567009  1.2875670086
6    -6.57048659 -1.074266908 -1.5090772378
.....

```

For the  $J$  observations, the variance is 6.116252 whereas for the  $S$  observations, the denominator is 9.118144. (with the expected ratio of 1.490806)

```

> pearson %>%
+   as.numeric %>%
+   matrix(nrow=85) %>%
+   round(4)

[,1]    [,2]    [,3]    [,4]    [,5]    [,6]
[1,] -0.7616  0.2194  0.3829 -0.2983  0.3597 -0.0790
[2,] -0.1450  0.1820 -0.1450 -0.5014  0.1567  0.2663
[3,] -0.8446  0.4634  0.1364 -0.1630 -0.2727  0.1660
[4,]  0.3745 -0.2795 -0.2795 -0.2658 -0.2658  0.6115
[5,]  1.2876 -0.6744 -0.6744  0.8935 -0.0935 -0.8612
[6,] -1.0743  1.8687 -0.7473 -0.0383  0.2908 -0.3673
.....

```

We can plot the residuals against the fitted values, to assess the assumption of constant variance.

```
# standardized residuals versus fitted values
plot(JS.roy1, resid(., type = "pearson") ~ fitted(.) ,
     abline = 0, id = 0.05)
```




Figure 2:

- The **Scale-Location** plot, also called Spread-Location or S-L plot, takes the square root of the absolute residuals in order to diminish skewness ( $\sqrt{|\text{residual}|}$ ) is much less skewed than  $|\text{residual}|$  for Gaussian zero-mean  $E$ ).
- The **Residual-Leverage** plot shows contours of equal Cook's distance, for values of `cook.levels` (by default 0.5 and 1) and omits cases with leverage one with a warning. If the leverages are constant (as is typically the case in a balanced aov situation) the plot uses factor level combinations instead of the leverages for the x-axis. (The factor levels are ordered by mean fitted value.)

```
par(mfrow=c(4,1))
plot(fittedmodel)
par(opar)
```

### 0.1.1 Residual Plots

A residual plot is a graph that shows the residuals on the vertical axis and the independent variable on the horizontal axis. If the points in a residual plot are randomly dispersed around the horizontal axis, a linear regression model is appropriate for the data; otherwise, a non-linear model is more appropriate.

Below the table on the left shows inputs and outputs from a simple linear regression analysis, and the chart on the right displays the residual (e) and independent variable (X) as a residual plot.

x	&	60	&	70	&	80	&	85	&	95	\\	\\hline
y	&	70	&	65	&	70	&	95	&	85	\\	\\hline
y.hat		65.411		71.849		78.288		81.507		87.945		
e		4.589		-6.849		-8.288		13.493		-2.945		

The residual plot shows a fairly random pattern - the first residual is positive, the next two are negative, the fourth is positive, and the last residual is negative. This random pattern indicates that a linear model provides a decent fit to the data.

Below, the residual plots show three typical patterns. The first plot shows a random pattern, indicating a good fit for a linear model. The other plot patterns are non-random (U-shaped and inverted U), suggesting a better fit for a non-linear model.